

1.

Ten thousand iterations:

$$\begin{aligned}E_{in} &= 0.5847 \\ \text{TestError} &= 0.3172\end{aligned}$$

One hundred thousand iterations:

$$\begin{aligned}E_{in} &= 0.4937 \\ \text{TestError} &= 0.2069\end{aligned}$$

One million iterations:

$$\begin{aligned}E_{in} &= 0.4354 \\ \text{TestError} &= 0.1310\end{aligned}$$

It seems like our hypothesis set H contain target function since Test Error is smaller as iteration time increase, but H looks like too big to achieve target function in one million iterations, under current learning rate.

Using *glmfit*

$$\text{TestError} = 0.1103$$

glmfit's running time is as fast as ten thousand iterations in previous method, and achieved 0.1103 Test Error, better than the best in previous.

Scale data result

$$\eta = 6.5, \text{ iteration} = 28, E_{in} = 0.4074$$

2 LFD 2.22

We know that:

$$E_D[E_{out}(g^{(D)})] = E_x[E_D[(g^{(D)}(x) - \bar{g}(x))^2] + (\bar{g}(x) - y(x))^2]$$

Then replace $y(x)$ with $f(x) + \epsilon$:

$$\begin{aligned} E_D[E_{out}(g^{(D)})] &= E_x[E_D[(g^{(D)}(x) - \bar{g}(x))^2] + (\bar{g}(x) - f(x))^2 + \epsilon^2 - 2\epsilon(\bar{g}(x) - f(x))] \\ &= E_x[var(x) + bias(x) + \epsilon^2 - 2\epsilon(\bar{g}(x) - f(x))] \\ &= var + bias + E_x[\epsilon^2 - 2\epsilon(\bar{g}(x) - f(x))] \\ &= var + bias + E_\epsilon[E_x[\epsilon^2 - 2\epsilon(\bar{g}(x) - f(x))|\epsilon]] \\ &= var + bias + E_\epsilon[\epsilon^2 - 2\epsilon(\bar{g}(x) - f(x))] \\ &= var + bias + E_\epsilon[\epsilon^2] - 2E_\epsilon[\epsilon(\bar{g}(x) - f(x))] \\ &= var + bias + var(\epsilon) - (E_\epsilon(\epsilon))^2 - 2E_\epsilon[\epsilon](\bar{g}(x) - f(x)) \end{aligned}$$

Since $E_\epsilon[\epsilon] = 0$

$$E_D[E_{out}(g^{(D)})] = var + bias + \sigma^2$$

Proved.

3 LFD 2.24

(a)

$$\begin{aligned} a &= \frac{y_2 - y_1}{x_2 - x_1} \\ b &= y_1 - \frac{y_2 - y_1}{x_2 - x_1}x_1 \\ &= \frac{y_1x_2 + y_2x_1}{x_2 - x_1} \\ \bar{g}(x) &= E_D\left[\frac{x_2^2 - x_1^2}{x_2 - x_1}x + \frac{x_1^2x_2 - x_2^2x_1}{x_2 - x_1}\right] \end{aligned}$$

Since x_1, x_2 's density function is $\frac{1}{2}, \frac{1}{2}$, respectively, then:

$$\begin{aligned}\bar{g}(x) &= \frac{1}{4} \int_{-1}^1 \int_{-1}^1 x_1 + x_2 dx_1 dx_2 + \frac{1}{4} \int_{-1}^1 \int_{-1}^1 x_1 x_2 dx_1 dx_2 \\ &= 0\end{aligned}$$

(b)

For n times, we generate x_1, x_2 from $[-1, 1]$, calculate its $g^{D_n}(x)$ and save it to a vector, after n times, we calculate $\bar{g}(x) = \frac{1}{n} \sum_{i=1}^n (g^{D_i})$, calculate $var(x) = \frac{1}{n} \sum_{i=1}^n (g^{D_i}(x) - \bar{g}(x))^2$, calculate $bias(x) = (\bar{g}(x) - f(x))^2$, calculate $E[E_{out}(g^D(x))] = \frac{1}{n} \sum_{i=1}^n (g^{D_i}(x) - f(x))^2$, then we generate total m points from $U[-1, 1]$, plug in those function and calculate mean value of $var(x)$, $bias(x)$, $E[E_{out}(g^D(x))]$ as var , $bias$, and $E[E_{out}(g^D)]$.

(c)

$E[E_{out}] = bias + var$, based on the experiment result.

$$E[E_{out}] = 0.5251$$

$$bias = 0.1963$$

$$var = 0.3289$$




Figure 1: $\bar{g}(x)$ and $f(x)$

(d)

$$\begin{aligned}
 E_{out} &= E[g(x) - f(x)] = E[(ax + b - x^2)^2] \\
 &= E[x^4] - 2aE[x^3] + (a^2 - 2b)E[x^2] + 2abE[X] + b^2 \\
 &= \frac{1}{2} \int_{-1}^1 x^4 dx - 2a \times \frac{1}{2} \int_{-1}^1 x^3 dx + \frac{1}{2}(a^2 - 2b) \int_{-1}^1 x^2 dx + ab \int_{-1}^1 x dx + b^2 \\
 &= \frac{1}{5} + \frac{1}{3}(a^2 - 2b) + b^2
 \end{aligned}$$

Then, take expectation of E_{out} with a , b replace by x_1 , x_2 in the (a):

$$\begin{aligned}
E[E_{out}] &= \frac{1}{5} + \frac{1}{3}E[(x_1 + x_2)^2 + 2x_1x_2] + E[x_1^2x_2^2] \\
&= \frac{1}{3} \frac{1}{4} \int_{-1}^1 \int_{-1}^1 (x_1 + x_2)^2 + 2x_1x_2 dx_1 dx_2 + \frac{1}{4} \int_{-1}^1 \int_{-1}^1 x_1^2 x_2^2 dx_1 dx_2 \\
&= \frac{8}{15}
\end{aligned}$$

For *bias*:

$$\begin{aligned}
bias &= E[bias(x)] = E[(\bar{g}(x) - f(x))^2] = E[x^4] \\
&= \frac{1}{2} \int_{-1}^1 x^4 dx = \frac{1}{5}
\end{aligned}$$

For *var*(*x*):

$$\begin{aligned}
var(x) &= E[(g(x) - \bar{g}(x))^2] = E[a^2x^2 + 2abx + E[b^2]] \\
&= E[a^2]x^2 + 2E[ab]x + b^2 \\
&= \frac{x^2}{4} \int_{-1}^1 \int_{-1}^1 (x_1 + x_2)^2 dx_1 dx_2 + 2 \times \frac{x}{4} \int_{-1}^1 \int_{-1}^1 (x_1 + x_2)(-x_1x_2) dx_1 dx_2 + \frac{x}{4} \int_{-1}^1 \int_{-1}^1 x_1^2 x_2^2 dx_1 dx_2 \\
&= \frac{2}{3}x^2 + \frac{1}{9}
\end{aligned}$$

Then *var* is :

$$\begin{aligned}
var &= E[var(x)] = E\left[\left(\frac{2}{3}x^2 + \frac{1}{9}\right)\right] \\
&= \frac{2}{3} \times \frac{1}{2} \int_{-1}^1 x^2 dx + \frac{1}{9} \\
&= \frac{1}{3}
\end{aligned}$$

4 LFD Exercise 3.4

(a)

$$\begin{aligned}\hat{y} &= Hy = HXw^* + H\epsilon \\ &= X(X^T X)^{-1}(X^T X)w^* + H\epsilon \\ &= Xw^* + H\epsilon\end{aligned}$$

Proved.

(b)

$$\begin{aligned}\hat{y} - y &= Xw^* + H\epsilon - Xw^* - \epsilon \\ &= (H - I)\epsilon\end{aligned}$$

So the matrix is $(H - I)$ where I is identity matrix with size of N .

(c)

$$\begin{aligned}E_{in}(w_{lin}) &= \frac{1}{N}(\hat{y} - y)^T(\hat{y} - y) \\ &= \frac{1}{N}(I - H)^2\epsilon^2\end{aligned}$$

By $(I - H)^k = I - H$:

$$E_{in}(w_{lin}) = \frac{1}{N}(I - H)\epsilon^2$$

5 LFD 3.4

(a)

$$e_n(w) = \begin{cases} 0 & , y_n w^T x_n > 1 \\ (1 - y_n w^T x_n)^2 & , y_n w^T x_n < 1 \end{cases}$$

It is continuous since:

$$\lim_{w: y_n w^T x_n \rightarrow 1} e_n(w) = 0$$

It is differentiable since:

$$\nabla e_n(w) = \begin{cases} 0 & , y_n w^T x_n > 1 \\ -2y_n(1 - y_n w^T x_n)x_n & , y_n w^T x_n < 1 \end{cases}$$

(b)

If $[sign(w^T x_n) \neq y_n] = 1$, then we know that $y_n w^T x_n \leq 0 < 1$, so in this case:

$$[sign(w^T x_n) \neq y_n] = 1 \leq e_n(w)$$

Similarly, if $[sign(w^T x_n) \neq y_n] = 0$, then $y_n w^T x_n \geq 0$, which means $e_n(w) = 0$ if $y_n w^T x_n < 1$; $e_n(w) > 1$ if $y_n w^T x_n > 1$, so in conclusion:

$$[sign(w^T x_n) \neq y_n] \leq e_n(w)$$

Hence, $\frac{1}{N} \sum_{n=1}^N e_n(w)$ is an upper bound for the in-sample classification error $E_{in}(w)$.

(c)

Use this bound for gradient descent is exactly Adaline algorithm.

If $y_n s_n = y_n w^T x_n \leq 1$, then update w :

$$w \leftarrow w - \eta \nabla e_n(w) = w - 2\eta y_n(1 - y_n w^T x_n)x_n = w - \eta'(y_n - w^T x_n)x_n$$

Where $s_n = w^T x_n$, and if $y_n s_n = y_n w^T x_n > 1$, do nothing, just as Adaline algorithm.

6 LFD 3.19

(a)

This transform will increase feature dimension to N , which will obviously slow down computation. At the same time, it will have $d_{vc} = N + 1$, which will have far more weaker generalization ability.

(b)

It seems good, but I worry that if γ is intensively small, feature itself will be smoked by very small γ .

(c)

It seems good.